

# Basics of Computing

Peter Ruprecht

[peter.ruprecht@colorado.edu](mailto:peter.ruprecht@colorado.edu)

[www.rc.colorado.edu](http://www.rc.colorado.edu)

Questions? #RC\_BasicSC

Link to survey on this topic: <http://tinyurl.com/rcpresurvey>

Slides:

[https://github.com/ResearchComputing/Final\\_Tutorials/tree/master/Basics\\_Supercomputing/2017\\_January](https://github.com/ResearchComputing/Final_Tutorials/tree/master/Basics_Supercomputing/2017_January)

# Outline

- Hierarchy: supercomputer to core
- Components of a compute node
- Interconnect
- Storage and filesystems
- Coprocessors
- Operating system

# Compute Hardware Architecture

Processing: Supercomputer – Node – Socket/CPU – Core

Memory: Distributed – RAM – L3 – L2 – L1

Interconnect network

Storage (disk)

Coprocessors

Non-cluster supercomputers (IBM Blue Gene, SGI UV)

Cloud?

*Parallelism at all levels of the computing system is the name of the game today ... SIMD, OpenMP, MPI*

# Clusters

## Stampede (TACC)



6400 nodes

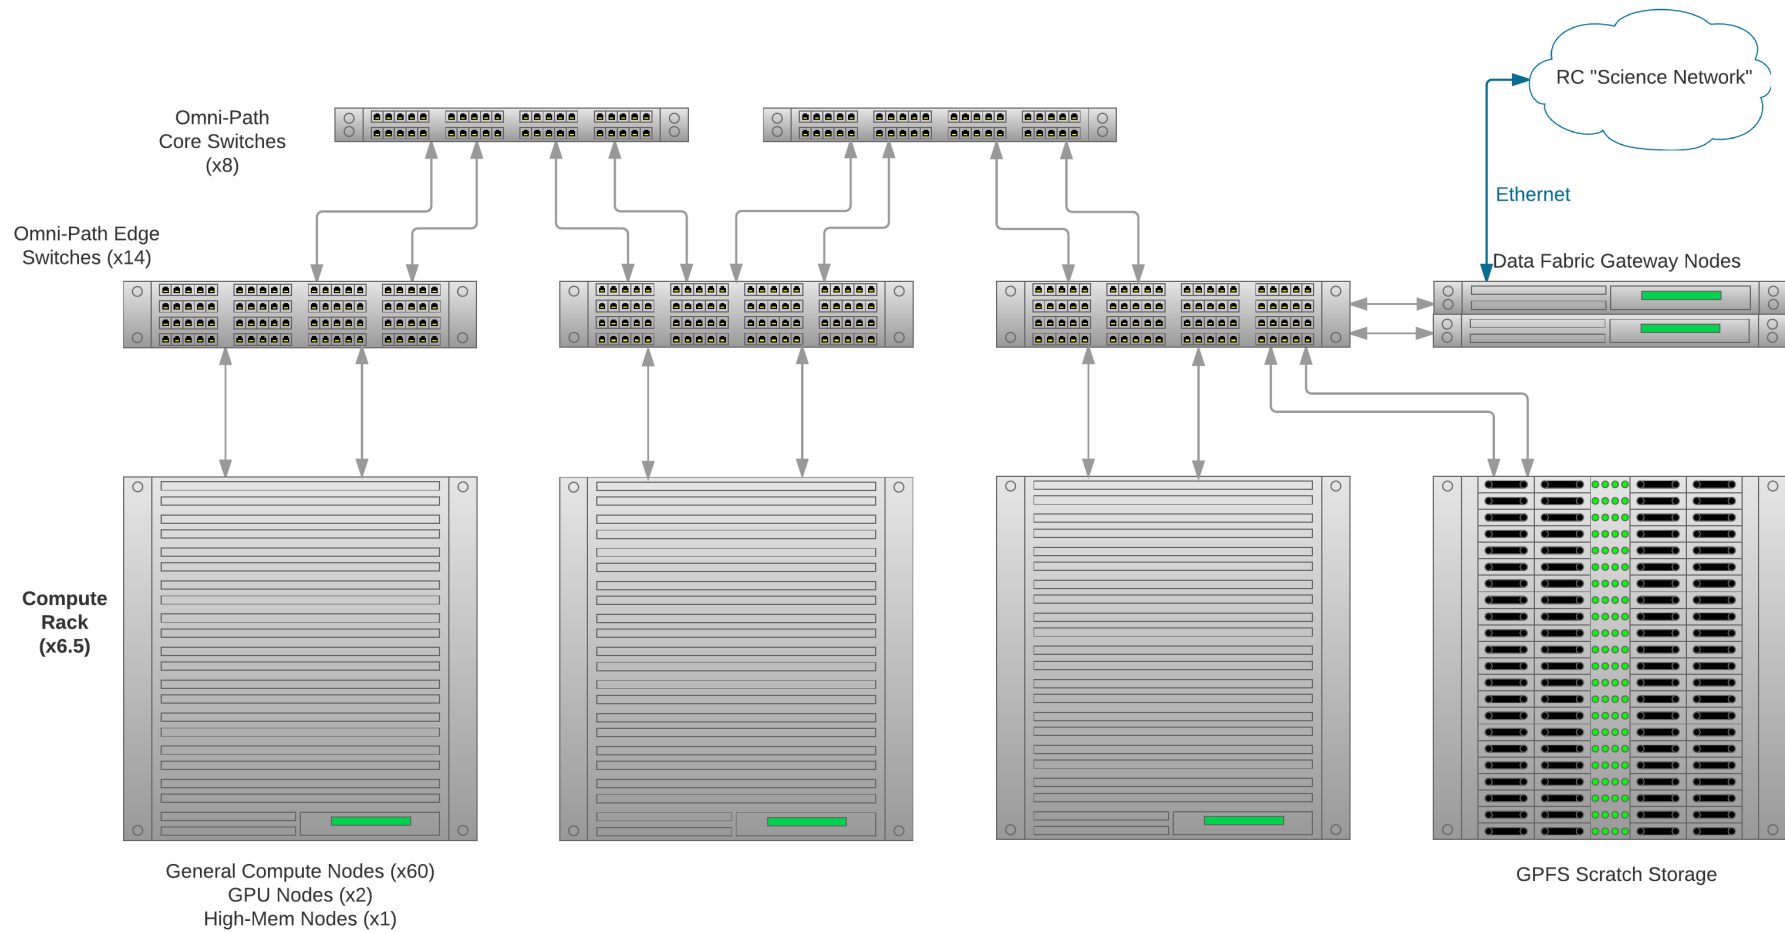
2 sockets per node, 8 cores per socket

56 Gbps Infiniband interconnect

14 PB parallel storage system


# SUMMIT SCHEMATIC

Peter Ruprecht | July 15, 2016




# Nodes and Processors

CPU sockets



Memory slots



Images: Supermicro, Intel

# Host Processor Types

- x86\_64 (Intel, AMD)
  - Most likely to encounter
  - General purpose
- Xeon Phi (Intel)
- POWER (IBM)
  - Blue Gene and future systems
- ARM
  - RISC, means fewer transistors
  - Low power

All processors operate at a certain frequency (several GHz) and can perform several instructions per cycle.

# Computer Bus Layout

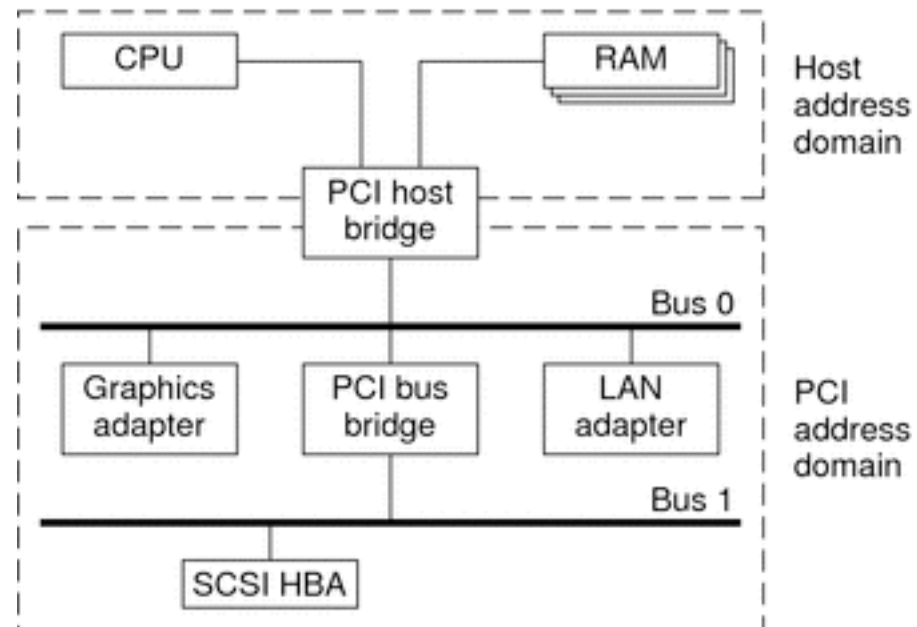


Image: Oracle



# Memory

- Holds data that is being calculated on, as well as computational instructions
- Closest memory to CPU is actually on the processor die
  - Cache: Level 1 and 2 are dedicated to a single core
  - Level 3 is larger and shared between all cores in a socket
- Next closest is RAM
- Can also access memory on other nodes, via RDMA (remote direct memory access)

*Shared memory* is local to one node and several process threads can share the same data addresses.

*Distributed memory* is on multiple nodes and each process normally has its own copy or part of the data.

# Interconnect

- How to access distributed memory across nodes?
- Need a fast, *low-latency* network.
- Current interconnect technologies
  - InfiniBand (40, 56, 100 Gbps)
  - Ethernet (10, 40, 100 Gbps)
  - Aries/Gemini (Cray)
  - OmniPath (100 Gbps)
- RDMA normally requires special interconnect-aware libraries; frequently these are included with MPI

Interconnect is what makes a bunch of nodes into a supercomputer!

# Storage

Different types of “disk” for different needs

- Local disk in the node, often SSD
- Shared scratch
  - Often accessed over cluster interconnect
  - Parallel filesystems, eg Lustre or GPFS
  - Traditionally tuned for high bandwidth, not high IOPS
  - May have a “burst buffer” layer in front of it
  - Short-term storage only!!
- Longer-term or archive
  - Often uses Hierarchical Storage Management

# Optimizing for Data Access

- Page Fault, file on IDE disk: 1,000,000,000 cycles
- Page Fault, file in buffer cache: 10,000 cycles
- Page Fault, file on ram disk: 5,000 cycles
- Page Fault, zero page: 3,000 cycles
- Main memory access: about 200 cycles
- L3 cache hit: about 52 cycles
- L1 cache hit: 2 cycles

**The Core i7 can issue 4 instructions per cycle. So a penalty of 2 cycles for L1 memory access means a missed opportunity for 7 instructions.**

# Co-Processors

Specialized, usually massively multi-core, separate from host processor (CPU).

- GPU (Graphics Processing Unit, “video card”)
  - Thousands of cores
  - Great for vectorizable or embarrassingly-parallel apps
  - Programming frameworks include CUDA, OpenACC, OpenCL
  - Many applications (eg, Matlab) support CUDA directly
- Xeon Phi (being phased out)
  - Dozens of cores
  - Existing x86 code can be directly recompiled to run on Phi
  - Function offload

# Operating Systems in HPC


Operating system: software that manages computer hardware and the processes running on it.

- MS Windows – hasn't really gotten a foothold
- Mac OSX – even less of a factor
- Unix / Linux – some variation will be on virtually any HPC system you use
  - May be a stripped-down version optimized for the particular hardware, or
  - May be a full distribution

# What is Linux?

- Part of the Unix family of operating systems.
- Started in early '90s by Linus Torvalds.
- Technically refers only to the kernel; software from the GNU project and elsewhere is layered on top to form a complete OS. Most is open source.
- Several full distributions are available – from enterprise-grade, like RHEL or SUSE, to more consumer-focused, like Ubuntu.
- Runs on everything from embedded systems to supercomputers.

# History of Linux



Brian Kernighan  
1970  
“space travel” to Unix



Dennis Ritchie  
1971  
C



Richard Stallman  
1983  
Gnu Not Unix



Linus Torvalds  
1991  
Linux kernel for personal computers



# Why use Linux?

- Linux command-line syntax may seem overwhelming to the new user, but:
- It's the default operating system on virtually all HPC systems
- It's extremely flexible
- It tries not to get in your way
- It's fast and powerful
- Open-source scientific applications are developed predominantly for Linux
- You can get started with a few basic commands and build from there

users

shell: bash, csh

programs      utilities

Linux kernel

Computer hardware

# References

1. Eijkhout, Victor. *Introduction to High-Performance Scientific Computing*, 2015.  
<https://bitbucket.org/VictorEijkhout/hpc-book-and-course/src>
2. Hager, G, and G Wellein. *Introduction to High Performance Computing for Scientists and Engineers*, CRC Press, 2010.
3. Levesque, John, and Gene Wagenbreth. *High Performance Computing*. CRC Press, 2010.
4. Neeman, Henry. *Supercomputing in Plain English*, High Performance Computing Workshop Series, <http://www.oscer.ou.edu/education.php>

# Questions?

- Email [rc-help@colorado.edu](mailto:rc-help@colorado.edu)
- Twitter: CUBoulderRC
- Link to survey on this topic:  
<http://tinyurl.com/curc-survey16>
- Slides:  
[https://github.com/ResearchComputing/Final\\_Tutorials/tree/master/Basics\\_Supercomputing/2017\\_January](https://github.com/ResearchComputing/Final_Tutorials/tree/master/Basics_Supercomputing/2017_January)